

# Not Hotdog Classifier Using fast.ai

This is a classic Silicon Valley: Season 4 Episode 4: Not Hotdog (HBO): <https://www.youtube.com/watch?v=pqTntG1RXY> (<https://www.youtube.com/watch?v=pqTntG1RXY>).

Blog: <https://medium.com/@timanglade/how-hbos-silicon-valley-built-not-hotdog-with-mobile-tensorflow-keras-react-native-ef03260747f3>  
(<https://medium.com/@timanglade/how-hbos-silicon-valley-built-not-hotdog-with-mobile-tensorflow-keras-react-native-ef03260747f3>).

The original code reference that I slightly modified:  
<https://github.com/pvhee/fastai-hotdog/blob/master/notebooks/hotdog.ipynb>  
(<https://github.com/pvhee/fastai-hotdog/blob/master/notebooks/hotdog.ipynb>).

Dataset downloaded from  
"<https://www.kaggle.com/dansbecker/hot-dog-not-hot-dog/downloads/hot-dog-not-hot-dog.zip>  
(<https://www.kaggle.com/dansbecker/hot-dog-not-hot-dog/downloads/hot-dog-not-hot-dog.zip>)."

**Just experimenting with my new EVGA RTX 2080 Ti XC Gaming GPU**

In [1]:

```
%reload_ext autoreload
%autoreload 2
%matplotlib inline
```

In [2]:

```
from fastai import *
from fastai.vision import *
import zipfile
```

## Run in specific GPU

In [3]:

```
# Update the GPU to run or delete this cell completely.
# I'm testing my new EVGA 2080 Ti XC Gaming gpu :)
import torch

for i in range(torch.cuda.device_count()):
    print(f"GPU NO. {i}\tNAME: {torch.cuda.get_device_name(i)}")

torch.cuda.set_device(0)
```

```
GPU NO. 0          NAME: GeForce RTX 2080 Ti
GPU NO. 1          NAME: GeForce GTX 1060 6GB
```

## Prepare the data

In [4]:

```
dataset = "hot-dog-not-hot-dog.zip"
data_path = Path("data-images/")
data_path
```

Out[4]:

```
PosixPath('data-images')
```

In [5]:

```
zip_ref = zipfile.ZipFile(dataset, "r")
zip_ref.extractall(data_path)
zip_ref.close()
```

In [6]:

```
np.random.seed(42)

bs = 64 # lower the batch size if your GPU has low memory or you are just using CPU. This network uses ResNet50
data = (ImageDataBunch
        .from_folder(data_path, train=".",
        valid_pct=0.2, ds_tfms=get_transforms(), size=224, bs=bs)
        .normalize(imagenet_stats))
```

In [7]:

```
data.show_batch(rows=3, figsize=(10,9))
```



In [8]:

```
data.classes, data.c, data
```

Out[8]:

```
(['hot_dog', 'not_hot_dog'], 2,  
ImageDataBunch;
```

```
Train: LabelList (1597 items)  
x: ImageList  
Image (3, 224, 224),Image (3, 2  
24, 224),Image (3, 224, 224),Ima  
ge (3, 224, 224),Image (3, 224,  
224)  
y: CategoryList  
not_hot_dog,not_hot_dog,not_hot  
_dog,not_hot_dog,not_hot_dog  
Path: data-images;
```

```
Valid: LabelList (399 items)  
x: ImageList  
Image (3, 224, 224),Image (3, 2  
24, 224),Image (3, 224, 224),Ima  
ge (3, 224, 224),Image (3, 224,  
224)  
y: CategoryList  
hot_dog,not_hot_dog,hot_dog,hot  
_dog,hot_dog  
Path: data-images;
```

```
Test: None)
```

# SOURCE CODE FOR TORCHVISION.MODELS.RESNET

See:

[https://pytorch.org/docs/stable/\\_modules/torchvision/models/resnet.html](https://pytorch.org/docs/stable/_modules/torchvision/models/resnet.html)  
([https://pytorch.org/docs/stable/\\_modules/torchvision/models/resnet.html](https://pytorch.org/docs/stable/_modules/torchvision/models/resnet.html))

This sample uses the resnet50 model. If your gpu or you are using a cpu to train, you can use resnet34, resnet18 or lower.

Good read about ResNet:

<https://towardsdatascience.com/introduction-to-resnets-c0a830a288a4>  
(<https://towardsdatascience.com/introduction-to-resnets-c0a830a288a4>)



## Now train em....

In [9]:

```
learn = cnn_learner(data, models.resnet50, metrics=accuracy)
```

In [10]:

```
learn.fit_one_cycle(5) # Experimented 5 epochs which resulted with high accuracy. Just preventing overfitting.
```

epoch	train_loss	valid_loss	accuracy	time
0	0.588014	0.448769	0.902256	00:00:00
1	0.425513	0.203998	0.944862	00:00:00
2	0.300958	0.086315	0.972431	00:00:00
3	0.210306	0.074175	0.969925	00:00:00
4	0.162175	0.055461	0.977444	00:00:00



In [11]:

```
learn.save("hotdog-stage-1") # saving the weights
learn.load("hotdog-stage-1");
```

In [12]:

```
interp = ClassificationInterpretation.from_learner(learn)
interp.plot_top_losses(9, figsize=(12,12), heatmap=False) # Some imperfections in prediction :)
```

Prediction/Actual/Loss/Probability

not\_hot\_dog/hot\_dog / 2.04 / 0.13



not\_hot\_dog/hot\_dog / 1.69 / 0.18



not\_hot\_dog/hot\_dog / 1.69 / 0.18



not\_hot\_dog/hot\_dog / 1.62 / 0.20



not\_hot\_dog/hot\_dog / 1.62 / 0.20




hot\_dog/not\_hot\_dog / 1.41 / 0.24



hot\_dog/not\_hot\_dog / 1.00 / 0.37



not\_hot\_dog/hot\_dog / 0.84 / 0.43



hot\_dog/not\_hot\_dog / 0.71 / 0.49



**Let's predict some HOTDOGS!!!**



## Prediction No. 1: Hotdog

In [13]:

```
hotdog1 = open_image("./predict-images/hotdog_1.jpeg")  
hotdog1
```

Out[13]:



In [14]:

```
pred_class, pred_idx, outputs = learn.predict(hotdog1)  
pred_class, pred_idx, outputs
```

Out[14]:

```
(Category hot_dog, tensor(0), tensor([1.0000e+00, 1.2579e-08]))
```

## Prediction No. 2: Not Hotdog

In [15]:

```
not_hotdog1 = open_image("./predict-images/not_hotdog_1.jpeg")  
not_hotdog1
```

Out[15]:



In [16]:

```
pred_class, pred_idx, outputs = learn.predict  
t(not_hotdog1)  
pred_class, pred_idx, outputs
```

Out[16]:

```
(Category not_hot_dog, tensor  
(1), tensor([0.3231, 0.6769]))
```

**Prediction No. 3: Hotdog**

In [17]:

```
hotdog2 = open_image("./predict-images/hotdog_2.jpeg")  
hotdog2
```

Out[17]:



In [18]:

```
learn.predict(hotdog2)
```

Out[18]:

```
(Category hot_dog, tensor(0), tensor([0.9972, 0.0028]))
```

**Prediction No. 4: Not Hotdog**

In [19]:

```
not_hotdog2 = open_image("./predict-images/not_hotdog_2.jpeg")  
not_hotdog2
```

Out[19]:



In [20]:

```
learn.predict(not_hotdog2)
```

Out[20]:

```
(Category not_hot_dog, tensor  
(1), tensor([0.0363, 0.9637]))
```

In [ ]: